

Code for project 1

October 9, 2018

1 Project 1 - code

It is all meant to be run in order

1.1 A whole heap of setup

1.1.1 Imports

These appear to be the standard imports

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from matplotlib import cm
from matplotlib.ticker import LinearLocator, FormatStrFormatter
from random import random, seed
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import cross_val_score
from imageio import imread
```

1.1.2 The Franke Function

We define the Franke Function, and do it plot of it. All of that is from the project description.

The only non-standard thing we do, is define a function `Franke_on_row`, which is just a wrapper to apply the Franke Function to a list with two elements.

```
In [2]: fig = plt.figure()
ax = fig.gca(projection='3d')

# Make data.
fx = np.arange(0, 1, 0.05)
fy = np.arange(0, 1, 0.05)
fx, fy = np.meshgrid(fx,fy)

def FrankeFunction(x,y):
```

```

term1 = 0.75*np.exp(-(0.25*(9*x-2)**2) - 0.25*((9*y-2)**2))
term2 = 0.75*np.exp(-((9*x+1)**2)/49.0 - 0.1*(9*y+1))
term3 = 0.5*np.exp(-(9*x-7)**2/4.0 - 0.25*((9*y-3)**2))
term4 = -0.2*np.exp(-(9*x-4)**2 - (9*y-7)**2)
return term1 + term2 + term3 + term4

# A wrapper to let us apply FrankeFunction to a list of the form row = [x, y]
def Franke_on_row(row):
    return FrankeFunction(row[0], row[1])

fz = FrankeFunction(fx, fy)


# Plot the surface.
surf = ax.plot_surface(fx, fy, fz, cmap=cm.coolwarm, linewidth=0, antialiased=False)

# Customize the z axis.
ax.set_zlim(-0.10, 1.40)
ax.zaxis.set_major_locator(LinearLocator(10))
ax.zaxis.set_major_formatter(FormatStrFormatter('%.02f'))

# Add a color bar which maps values to colors.
fig.colorbar(surf, shrink=0.5, aspect=5)

plt.show()

```



1.1.3 Regression methods

We define the classes `polynomialOLS`, `polynomialRidge`, and `polynomialLasso`. Each will take a degree, and the later two also a parameter `lamb`. They each implement `train` and `predict` functions. The idea is that one initiates a polynomial method `pm` with a given degree, then we can simply pass a two-column matrix `XY` with `x` and `y` values to `train` and `predict`, the classes themselves will then call `polymatrix` to get a matrix with all the terms needed to get a polynomial of the given degree. The Ridge regression class also takes care to center the input. That is all just to make calling the functions easier.

Each class also has `plotfunc`, `coefficients` and `info` methods.

The `plotfunc` is simply a wrapper used to predict on a single pair of points `(x,y)`. I use it because I had problems with applying `predict` when plotting.

The `coefficients` methods returns a column vector consisting of the beta functions.

The `info` method does some printing we use later to make nice output.

```
In [3]: # We are only going to work with input in two dimensions
        # and output in one.
        # Hence all our code will be writting with this assumption

        # Given a matrix XY with two columns x,y return an array with columns x, y, x^2, xy, y^2
def polymatrix(XY, n):
    X = XY[:,[0]]
    Y = XY[:,[1]]
    ones = np.ones((len(X), 1))
    pl = np.concatenate([(np.power(X,k-i)*np.power(Y,i)) for k in range(1, n+1) for i in range(0, k)])
    return np.hstack((ones, pl))

# This is a class for doing the most basic OLS regression
class myOLS :

    # We have one class virable beta which has the coefficients for the linear regression
    def __init__(self):
        self.beta = 0

    # This method sets beta on the given training matrix xb with know values y.
    def train(self, xb, y):
        self.beta = np.linalg.inv(xb.T.dot(xb)).dot(xb.T).dot(y)

    # Use our beta to predict on a (correctly formatted) matrix xb.
    def predict(self, xb):
        return xb.dot(self.beta)

# This class is simply a wrapper for the myOLS class
# It simply takes care of transforming the raw input into proper polynomial input
class polynomialOLS :

    def __init__(self, n):
```

```

        self.polyOLS = myOLS()
        self.degree = n

# Train our polynomial OLS.
# We take matrix XY with two columns x,y and given zs as output.
# First we use polymatrix to transform XY to matrix with rows 1, x, y, x^2, xy, y^2
# Then we feed that as the training data to an ordinary OLS
def train(self, XY, z):
    xb = polymatrix(XY, self.degree)
    self.polyOLS.train(xb,z)

# Predict our polynomial solution
# As in train we take a matrix XY with two columns x,y.
# First we use polymatrix to transform XY to matrix with rows 1, x, y, x^2, xy, y^2
# Then use that as input to predict on an ordinart OLS
def predict(self, XY):
    xb = polymatrix(XY, self.degree)
    return self.polyOLS.predict(xb)

# This is only used for plotting.
# It is just a wrapper function to predict the value a single (x,y) pair
# Given two numbers x,y it just applies predict to an array with x and y as columns
def plotfunc(self, x, y):
    return self.predict(np.array([[x, y]]))

# Prints the coefficients beta
def coefficients(self):
    return self.polyOLS.beta

# A function to print some pretty information about this regression model
def info(self):
    print("OLS regression, trying to fit a polynomial of degree %d " % self.degree)

class polynomialRidge :

    def __init__(self, n, lamb):
        self.degree = n
        self.lamb = lamb
        self.beta = 0
        self.intercept = 0

# takes a column or row vector x, and returns x - bar{x},
# where bar{x} is the avarage of the entries in x
# Also works on any size matrix, but that is not what we have in mind
def center(self, x) :
    return x - np.mean(x)

# Train our polynomial ridge regression.

```

```

# We take matrix XY with two columns x,y and given zs as output.
# First we transform XY to matrix with columns x, y, x^2, xy, y^2 ...
# Then we center each row.
# Finally we compute beta according to formula from Hastings (3.44)
def train(self, XY, z):
    xb = polymatrix(XY, self.degree)
    xb = np.delete(xb, 0, 1) # removes the column of 1s
    xb_centered = np.apply_along_axis(self.center, 0, xb) # center each column
    xbrows = len(xb[0,:]) # number of rows in xb
    b0 = self.center(z)
    self.beta = (np.linalg.inv(xb_centered.T.dot(xb_centered) + self.lamb*np.identity(xbrows))
    self.intercept = np.mean(z) - (np.mean(xb, axis = 0)).dot(self.beta)

# Predict our polynomial solution
# As in train we take a matrix XY with two columns x,y.
# First we transform XY to matrix with columns x, y, x^2, xy, y^2 ...
# Then use beta and intercept to predict the value
def predict(self, XY):
    xb = polymatrix(XY, self.degree)
    xb = np.delete(xb, 0, 1) # removes the column of 1s
    return xb.dot(self.beta) + self.intercept

# This is only used for plotting.
# It is just a wrapper function to predict the value a single (x,y) pair
# Given two numbers x,y it just applies predict to an array with x and y as columns
def plotfunc(self, x, y):
    return self.predict(np.array([[x, y]]))

# Returns the coefficients beta (a long with intercept) as a column vector
# vstack puts the intercept on top of the other betas
def coefficients(self):
    return np.vstack((self.intercept, self.beta))

# A function to print some pretty information about this regression model
def info(self):
    print("Ridge regression, trying to fit a polynomial of degree %d " % self.degree)
    print("Lambda = %.4f" % self.lamb)

class polynomialLasso :

    def __init__(self, n, lamb):
        self.lasso = Lasso(alpha = lamb)
        self.degree = n
        self.lamb = lamb

# Train our polynomial lasso.
# We take matrix XY with two columns x,y and given zs as output.
# First we transform XY to matrix with columns x, y, x^2, xy, y^2 ...

```

```

# Then we center each row.
# Finally we compute beta according to formula from Hastings (3.44)
def train(self, XY, z):
    xb = polymatrix(XY, self.degree)
    xb = np.delete(xb, 0, 1) # removes the column of 1s
    self.lasso.fit(xb, z)

# Predict our polynomial solution
# As in train we take a matrix XY with two columns x,y.
# First we transform XY to matrix with columns x, y, x^2, xy, y^2 ...
# Then use beta and intercept to predict the value
def predict(self, XY):
    xb = polymatrix(XY, self.degree)
    xb = np.delete(xb, 0, 1) # removes the column of 1s
    return np.c_[self.lasso.predict(xb)]

# This is only used for plotting.
# It is just a wrapper function to predict the value a single (x,y) pair
# Given two numbers x,y it just applies predict to an array with x and y as columns
def plotfunc(self, x, y):
    return self.predict(np.array([[x, y]]))

# Returns the coefficients (including intercept) as a column vector
# reshape(-1,1) makes the array of coef_ into a column vector
# we put intercept_ on top of it.
def coefficients(self):
    # resha
    return np.vstack((self.lasso.intercept_, self.lasso.coef_.reshape(-1,1)))

def info(self):
    print("Lasso regression, trying to fit a polynomial of degree %d " % self.degree)
    print("Lambda = %.4f" % self.lamb)

```

1.1.4 Test of regression functions

We verify that the classes we build yield the same outcome as the build-in functions of sci-kit learn.

```

In [4]: print("Test of regression methods and polymatrix")
        print("-----")
        print("")
        print("Testing polymatrix. Making first a random (4,2) array")
        test_XY = np.random.rand(100,2)
        test_z = np.c_[np.apply_along_axis(Franke_on_row, 1, test_XY)]
        #print(test_XY)
        print("")

        for test_degree in range(1,6) :

```

```

print("For degree %d the mean difference between polmatrix and PolynomialFeatures are")
poly = PolynomialFeatures(test_degree)
print(np.mean(polmatrix(test_XY, test_degree) - poly.fit_transform(test_XY)))
print("")

print("-----")
print("")
print("Testing how my ols and ridge methods differ from the sci-kit learn ones")
print("")
for test_degree in range(1,6) :

    print("We look at a polynomial of degree %d" % test_degree)
    poly = PolynomialFeatures(test_degree)
    poly_XY = poly.fit_transform(test_XY)
    # We remove the leading column of 1s
    # This is done since we call the fit functions without setting fit_intercept to false
    # So in essence, the fit of OLS and more importantly Ridge makes sure to find the intercept
    poly_XY = np.delete(poly_XY, 0, 1)

    # Setup "my" OLS regression
    myols = polynomialOLS(test_degree)
    myols.train(test_XY, test_z)
    # Setup sci-kit learns OLS regression
    scikit_ols = LinearRegression()
    scikit_ols.fit(poly_XY, test_z)
    # make a column vector of betas for scikit_ols
    skols_betas = np.vstack((scikit_ols.intercept_, scikit_ols.coef_.reshape(-1,1)))
    print("Mean difference of hand computed intercept and coeffs for OLS")
    print(np.mean(myols.coefficients() - skols_betas))

    # Setup "my" ridge regression
    myridge = polynomialRidge(test_degree, 0.5)
    myridge.train(test_XY, test_z)
    # Setup sci-kit learns ridge regression
    scikit_ridge = Ridge(alpha = 0.5)
    scikit_ridge.fit(poly_XY, test_z)
    skridge_betas = np.vstack((scikit_ridge.intercept_, scikit_ridge.coef_.reshape(-1,1)))
    print("Mean difference of hand computed intercept and coeffs for Ridge")
    print(np.mean(myridge.coefficients() - skridge_betas))

```

Test of regression methods and polmatrix

Testing polmatrix. Making first a random (4,2) array

For degree 1 the mean difference between polmatrix and PolynomialFeatures are

0.0

For degree 2 the mean difference between polmatrix and PolynomialFeatures are
0.0

For degree 3 the mean difference between polmatrix and PolynomialFeatures are
2.524344530066211e-19

For degree 4 the mean difference between polmatrix and PolynomialFeatures are
3.493005673427315e-19

For degree 5 the mean difference between polmatrix and PolynomialFeatures are
2.3215042611641024e-19

Testing how my ols and ridge methods differ from the sci-kit learn ones

We look at a polynomial of degree 1

Mean difference of hand computed intercept and coefs for OLS

5.181040781584064e-16

Mean difference of hand computed intercept and coefs for Ridge

3.700743415417188e-17

We look at a polynomial of degree 2

Mean difference of hand computed intercept and coefs for OLS

-5.782411586589357e-16

Mean difference of hand computed intercept and coefs for Ridge

7.517135062566164e-17

We look at a polynomial of degree 3

Mean difference of hand computed intercept and coefs for OLS

5.004885395010206e-14

Mean difference of hand computed intercept and coefs for Ridge

-4.163336342344337e-18

We look at a polynomial of degree 4

Mean difference of hand computed intercept and coefs for OLS

-1.777851939740079e-12

Mean difference of hand computed intercept and coefs for Ridge

-3.7007434154171884e-18

We look at a polynomial of degree 5

Mean difference of hand computed intercept and coefs for OLS

3.660379605239345e-11

Mean difference of hand computed intercept and coefs for Ridge

-5.683284530819253e-17

1.1.5 Cross validation

```
In [5]: # This function does k-fold cross validation
# It takes a k, a matrix XY with two columns of inputs, and corresponding true values z
# It also takes a model which we assume have train, predict, and coefficients functions.
# Both should work on these matrices
# We return the expected error, and a list of computed coefficients for our model
def cross_validation(k, XY, z, model):

    sqdiffs = []
    clength = len(XY[:,0]) # the length of the first column in XY
    permutation = np.random.permutation(clength) # a permutation of the indexes of rows
    partitions = np.array_split(permutation, k) # The permutation is divided in to k alms

    for i in range(0,k) :
        # create a mask to pick everything but the elements in the i'th partition
        mask = np.ones(clength,dtype=bool)
        mask[partitions[i]] = 0
        # now train on everything but the i'th partition
        model.train(XY[mask, :], z[mask, :])
        # make the mask the picks only the elements of the i'th partition
        notmask = np.invert(mask)
        # update the sqdiffs of predicted values and the true values in the i'th partition
        zpredict = model.predict(XY[notmask, :])
        sqdiffs.append(np.power(zpredict - z[notmask, :],2))

    pred_mse = np.mean(sqdiffs)
    return pred_mse
```

1.1.6 Statistics functions

Just a small collection of statistics functions we will use to evaluate the preformance of our models

```
In [6]: def MSE(x, y):
    return np.mean(np.power(x-y, 2))

def r2d2score(ytrue, ypredict):
    truemean = np.mean(ytrue)
    return 1 - (np.sum(np.power(ytrue-ypredict, 2))/np.sum(np.power(ytrue-truemean, 2)))

# variance of a column of row vector x
def variance(x):
    return np.mean(np.power(x - np.mean(x), 2))

# Given input data XY, z, and a model.
# We first train the model on XY, z.
# Then we compute the MSE and R2 score of the prediction vs the true values
# After thinking it over, this just seems boring, so it wont be used
def basic_tests(XY, z, model) :
```

```

model.train(XY, z)
zpredict = model.predict(XY)
mse = MSE(z, zpredict)
r2 = r2d2score(z, zpredict)
model.info()
print("MSE = %.4f" % mse)
print("R2 score = %.4f" %r2)
print("-----")

# Given a model already trained on some data, and some new data XY, z
# Compute MSE and R2 scores of the models prediction of the newdata against it true new
def test_against_new(newXY, newz, model):
    zpredict = model.predict(newXY)
    mse = MSE(newz, zpredict)
    r2 = r2d2score(newz, zpredict)
    model.info()
    print("MSE = %f" % mse)
    print("R2 score = %f" % r2)
    print("-----")

```

1.1.7 Statistics functions test

```

In [7]: print("Testing the basic statistics functions")
        print("")
        # make two lists of random numbers
        # we need the ravel or the build in r2_score behaves in an odd way
        xs = np.random.rand(1,100).ravel()
        ys = np.random.rand(1,100).ravel()
        mse_diff = MSE(xs,ys) - mean_squared_error(xs,ys)
        print("MSE difference between mine and sci-kit learns: %f " % mse_diff)
        r2_diff = r2d2score(xs,ys) - r2_score(xs,ys)
        print("R2 difference between mine and sci-kit learns: %f " % r2_diff)

```

Testing the basic statistics functions

```

MSE difference between mine and sci-kit learns: 0.000000
R2 difference between mine and sci-kit learns: 0.000000

```

1.1.8 BV estimate and confidence intervals

Functions to estimate the bias and variance and confidence intervals of a given model. They are just minor modifications of the confidence interval code, and really one should probably do both (or all 3) in one go.

```

In [8]: # This function tries to estimate the variance and bias of our model
        # It takes a k, a matrix XY with two columns of inputs, and corresponding true values z

```

```

# It also takes a model which we assume have train, predict, and coefficients functions.
# We return the variance and bias of computing the model on k subsets of XY
# when compared to a single "true" set
def BV_estimate(k, XY, z, model):

    clength = len(XY[:,0]) # the length of the first column in XY
    permutation = np.random.permutation(clength) # a permutation of the indexes of rows
    partitions = np.array_split(permutation, k+1) # The permutation is divided in to k+1
    preds_list = [] # this will be a list of the predicted out comes

    # create a mask to pick everything in the k'th partition
    # this is the partition we will compare against
    fixed_mask = np.zeros(clength,dtype=bool)
    fixed_mask[partitions[k]] = 1

    for i in range(0,k) :
        # create a mask to pick everything but the i'th partition
        mask = np.ones(clength,dtype=bool)
        mask[partitions[i]] = 0
        mask[partitions[k]] = 0
        # now train on the i'th partition
        model.train(XY[mask, :], z[mask, :])
        # now predict what will happen on the fixed set
        zpredict = model.predict(XY[fixed_mask, :])
        preds_list.append(zpredict)

    # transform from list of column vectors to a single matrix
    # each row contains the prediction of a single xy point
    preds_matrix = np.hstack(preds_list)

    vari = np.mean(np.apply_along_axis(variance, 1, preds_matrix))
    bias = np.mean(np.power(np.mean(preds_matrix - z[fixed_mask, :], axis=1), 2))
    #mse = np.mean(np.power(preds_matrix - z[fixed_mask, :], 2))

    return bias, vari

##usage example:
#pRidge = polynomialOLS(4)
#print("Using 9 folds to predict variance and bias for a polynomial of degree %i." % 2)
#print(BV_estimate(9, XY, z, pRidge))

# This function estimates the variance of the betas
# It takes a k, a matrix XY with two columns of inputs, and corresponding true values z
# It also takes a model which we assume have train, predict, and coefficients functions.
# We return the variance the the betas when computing the model on k subsets of XY
def beta_variance(k, XY, z, model):

    clength = len(XY[:,0]) # the length of the first column in XY

```

```

permutation = np.random.permutation(clength) # a permutation of the indexes of rows
partitions = np.array_split(permutation, k) # The permutation is divided in to k alms
beta_list = [] # this will be a list of the betas for each training session

for i in range(0,k) :
    # create a mask to pick everything but the elements in the i'th partition
    mask = np.ones(clength,dtype=bool)
    mask[partitions[i]] = 0
    # now train on everything but the i'th partition
    model.train(XY[mask, :], z[mask, :])
    beta_list.append(model.coefficients())

# transform from list of column vectors to a single matrix
# each row contains the "same beta" computed from different training session
beta_matrix = np.hstack(beta_list)

variances = np.apply_along_axis(variance, 1, beta_matrix)

return variances

##usage example
#pRidge = polynomialOLS(4)
#print("Using 10 folds to compute variances of betas for a polynomial of degree %i." % 2)
#print(beta_variance(10, XY, z, pRidge))

```

1.2 Acutally working with the models or solving parts a,b, c

1.2.1 Data for all our test

```

In [53]: # Setup input data.
# We make numberofpoints points
# the x's and y's are chosen randomly
# the z's according to the fomula and with some noise
number_of_points = 1000
noise = 0.2
XY = np.random.rand(number_of_points, 2) # a matrix of random numbers with 2 columns of

# we have to flip the normal dist array to get the right dimensions
z = np.c_[np.apply_along_axis(Franke_on_row, 1, XY)] + np.c_[noise*np.random.normal(0,1

# Now we make some "clean" data
# Models will never be trained on this, only tested against it
cleanXY = np.random.rand(100, 2)
cleanz = np.c_[np.apply_along_axis(Franke_on_row, 1, cleanXY)] + np.c_[noise*np.random.

```

1.2.2 A, OLS

```
In [54]: print("Testing OLS methods for polynomial of degrees 1,2,3,4, and 5.")
print("")
print("First we use 10 fold cross validation to decide which model is best.")
print("")

for n in range(1,6) :
    pOLS = polynomialOLS(n)
    pOLS.info()
    cv = cross_validation(10, XY, z, pOLS)
    print("The predicted error is %f" % cv)
    print("")

print("-----")
print("")
print("Testing 5 degree polynomial model against new data")

# When I ran this to lowest predicted error came from the polynomial of degree 5
# so we will go with that model

# We test the actual statistics against the unseen (clean) data
bestOLS = polynomialOLS(5)
bestOLS.train(XY, z)
print("The coefficients for our best OLS model is:")
print(bestOLS.coefficients())
test_against_new(cleanXY, cleanz, bestOLS)

# Now we compute the bias and variance of our model
# and estimate the variance of the parameters beta

print("")
print("Using 9 folds to predict variance and bias")
b, v = BV_estimate(9, XY, z, bestOLS)
print("The bias is %f" % b)
print("The variance is %f" % v)
print("")
print("Using 10 folds to compute variances of betas")
print(beta_variance(10, XY, z, bestOLS))
```

Testing OLS methods for polynomial of degrees 1,2,3,4, and 5.

First we use 10 fold cross validation to decide which model is best.

OLS regression, trying to fit a polynomial of degree 1
The predicted error is 0.062884

OLS regression, trying to fit a polynomial of degree 2
The predicted error is 0.058462

OLS regression, trying to fit a polynomial of degree 3
The predicted error is 0.049181

OLS regression, trying to fit a polynomial of degree 4
The predicted error is 0.046228

OLS regression, trying to fit a polynomial of degree 5
The predicted error is 0.044361

Testing 5 degree polynomial model against new data
The coefficients for our best OLS model is:

```
[[ 0.50487456]
 [ 6.84139196]
 [ 3.40447963]
 [-28.83003238]
 [-13.51823062]
 [-8.49877932]
 [ 34.9506793 ]
 [ 36.27449647]
 [ 26.08048924]
 [-9.02317215]
 [-8.29829456]
 [-47.10130772]
 [-3.18784026]
 [-41.11659721]
 [ 32.0059086 ]
 [-5.18824459]
 [ 19.65607878]
 [ 3.94561365]
 [-1.70731701]
 [ 21.06836119]
 [-18.13575605]]
```

OLS regression, trying to fit a polynomial of degree 5
MSE = 0.046539
R2 score = 0.636927

Using 9 folds to predict variance and bias
The bias is 0.043213
The variance is 0.000132

Using 10 folds to compute variances of betas
[1.3355253e-03 2.37302168e-01 1.69367710e-01 8.30320758e+00
1.61226390e+00 3.36708078e+00 4.22657606e+01 4.71033533e+00
1.22514595e+01 1.76142175e+01 4.01009946e+01 1.26645999e+01]

```
7.45089345e+00 1.43587585e+01 2.26653797e+01 5.09400433e+00
3.50051075e+00 1.36496341e+00 2.68100217e+00 2.12393538e+00
3.97723740e+00]
```

1.2.3 A - a plot

We do a plot of our predicted model - franke function. If this had been a great model we should have seen a very flat plot, we don't quite

```
In [55]: # retrain model as it has been modified by BV and variance computations
         bestOLS.train(XY, z)

         # Do a plot of the best model - franke
         fig = plt.figure()
         ax = fig.gca(projection='3d')


         # Make plot points .
         xpoints = np.arange(0, 1, 0.05)
         ypoints = np.arange(0, 1, 0.05)
         xm, ym = np.meshgrid(xpoints, ypoints)

         # Plot the real and predicted surfaces.
         diffsurf = ax.plot_surface(xm, ym, np.vectorize(bestOLS.plotfunc)(xm, ym) - FrankeFunction(xm, ym))

         # Customize the z axis.
         ax.set_zlim(-0.50, 0.5)
         ax.zaxis.set_major_locator(LinearLocator(10))
         ax.zaxis.set_major_formatter(FormatStrFormatter('%.02f'))

         # Add a color bar which maps values to colors.
         #fig.colorbar(realsurf, shrink=0.5, aspect=5)

         plt.show()
```



1.2.4 B, ridge regression - model selection

```
In [56]: print("Testing Ridge methods for polynomial of degrees 1,2,3,4, and 5.")
print("")
print("First we use 10 fold cross validation to decide which model is best.")
print("")

#Since printing out all the results of cross validation with different lambdas get tedious
# we make a list of cv values and corresponding degrees and lambdas.
# Then we can just sort by the cv value to pick the best model
cv_list = []

for n in range(1,6) :
    for a in [0.1, 0.2, 0.5, 1] :
        pRidge = polynomialRidge(n, a)
        cv = cross_validation(10, XY, z, pRidge)
        cv_list.append((n,a,cv))

cv_list.sort(key=lambda tup: tup[2]) #sort the list by the cv value

best_tup = cv_list[0]

print("The best model had a degree %d polynomial and a lambda of %f " % (best_tup[0], best_tup[1]))
print("It had a predicted error of %f" % best_tup[2])

# NOTE: Since this is not entirely deterministic, we some times get different answers.
# Mostly I get a degree 4 polynomial with a lambda of 0.1
```


Testing Ridge methods for polynomial of degrees 1,2,3,4, and 5.

First we use 10 fold cross validation to decide which model is best.

The best model had a degree 5 polynomial and a lambda of 0.100000
It had a predicted error of 0.049839

1.2.5 B - working with chosen model

```
In [57]: # We will work with a degree 5 polynomial and a lambda of 0.1
# We test the actual statistics against the unseen (clean) data
bestRidge = polynomialRidge(5, 0.1)
bestRidge.train(XY, z)
print("The coefficients for our best Ridge model is:")
print(bestRidge.coefficients())
test_against_new(cleanXY, cleanz, bestRidge)

# Now we compute the bias and variance of our model
# and estimate the variance of the parameters beta

print("")
print("Using 9 folds to predict variance and bias")
b, v = BV_estimate(9, XY, z, bestRidge)
print("The bias is %f" % b)
print("The variance is %f" % v)
print("")
print("Using 10 folds to compute variances of betas")
print(beta_variance(10, XY, z, bestRidge))
```

The coefficients for our best Ridge model is:

```
[[ 1.06612447]
 [-0.43069641]
 [ 0.10905763]
 [-1.5575589 ]
 [ 0.99764565]
 [-2.0988343 ]
 [ 0.70765274]
 [ 0.73417389]
 [-0.66134346]
 [-0.28112356]
 [ 0.89019451]
 [ 0.6225893 ]
 [-0.21970549]
 [-0.32754434]
 [ 0.87460214]
 [-0.6857009 ]
 [-0.22535477]]
```

```

[-0.26390879]
[-0.08609822]
[ 0.23636058]
[ 0.76602987]]
Ridge regression, trying to fit a polynomial of degree 5
Lambda = 0.1000
MSE = 0.058986
R2 score = 0.539831
-----

```

```

Using 9 folds to predict variance and bias
The bias is 0.054887
The variance is 0.000056

```

```

Using 10 folds to compute variances of betas
[0.00033641 0.00409355 0.00262426 0.00327892 0.00884094 0.00301148
 0.00212812 0.00403362 0.0044825  0.00108847 0.001235  0.00170429
 0.00141348 0.00185033 0.00048559 0.00107178 0.00369705 0.00535533
 0.00346365 0.00359964 0.00189097]

```

1.2.6 B - the plot

```

In [58]: # retrain model as it has been modified by BV and variance computations
         bestRidge.train(XY, z)

         # Do a plot of the best model - franke
         fig = plt.figure()
         ax = fig.gca(projection='3d')

         # Make plot points .
         xpoints = np.arange(0, 1, 0.05)
         ypoints = np.arange(0, 1, 0.05)
         xm, ym = np.meshgrid(xpoints, ypoints)


         # Plot the real and predicted surfaces.
         diffsurf = ax.plot_surface(xm, ym, np.vectorize(bestRidge.plotfunc)(xm,ym) - FrankeFunc)

         # Customize the z axis.
         ax.set_zlim(-0.50, 0.5)
         ax.zaxis.set_major_locator(LinearLocator(10))
         ax.zaxis.set_major_formatter(FormatStrFormatter('%.02f'))

         # Add a color bar which maps values to colors.
         #fig.colorbar(realsurf, shrink=0.5, aspect=5)

         plt.show()

```



1.2.7 C, lasso regression - model selection

```
In [59]: print("Testing Lasso methods for polynomial of degrees 1,2,3,4, and 5.")
print("")
print("First we use 10 fold cross validation to decide which model is best.")
print("")

#Since printing out all the results of cross validation with different lambdas get tedious
# we make a list of cv values and corresponding degrees and lambdas.
# Then we can just sort by the cv value to pick the best model
cv_list = []

for n in range(1,6) :
    for a in [0.1, 0.2, 0.5, 1] :
        pLasso = polynomialLasso(n, a)
        cv = cross_validation(10, XY, z, pLasso)
        cv_list.append((n,a,cv))

cv_list.sort(key=lambda tup: tup[2]) #sort the list by the cv value

best_tup = cv_list[0]

print("The best model had a degree %d polynomial and a lambda of %f " % (best_tup[0], best_tup[1]))
print("It had a predicted error of %f" % best_tup[2])

# NOTE: Since this is not entirely deterministic, we sometimes get different answers.
```

```
# In fact it seems incredibly unstable.
# I assume a big part of this is that the lasso models sets almost all betas to zero
# When I did it, I got a degree 2 polynomial with a lambda of 0.2
```

Testing Lasso methods for polynomial of degrees 1,2,3,4, and 5.

First we use 10 fold cross validation to decide which model is best.

The best model had a degree 4 polynomial and a lambda of 0.100000
It had a predicted error of 0.124234

1.2.8 C - working with chosen model

```
In [66]: # We will work with a degree 5 polynomial and a lambda of 0.1
# We test the actual statistics against the unseen (clean) data
bestLasso = polynomialLasso(4, 0.1)
bestLasso.train(XY, z)
print("The coefficients for our best Lasso model is:")
print(bestLasso.coefficients())
test_against_new(cleanXY, cleanz, bestLasso)

# Now we compute the bias and variance of our model
# and estimate the variance of the parameters beta

print("")
print("Using 9 folds to predict variance and bias")
b, v = BV_estimate(9, XY, z, bestLasso)
print("The bias is %f" % b)
print("The variance is %f" % v)
print("")
print("Using 10 folds to compute variances of betas")
print(beta_variance(10, XY, z, bestLasso))
```

The coefficients for our best Lasso model is:

```
[[ 0.39662301]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
 [-0.         ]
```

```

[-0.      ]
[-0.      ]]
Lasso regression, trying to fit a polynomial of degree 4
Lambda = 0.1000
MSE = 0.130758
R2 score = -0.020091
-----

Using 9 folds to predict variance and bias
The bias is 0.128723
The variance is 0.000007

Using 10 folds to compute variances of betas
[2.51423352e-05 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00]

```

1.2.9 C - the plot

```

In [61]: # retrain model as it has been modified by BV and variance computations
         bestLasso.train(XY, z)

         # Do a plot of the best model - franke
         fig = plt.figure()
         ax = fig.gca(projection='3d')

         # Make plot points .
         xpoints = np.arange(0, 1, 0.05)
         ypoints = np.arange(0, 1, 0.05)
         xm, ym = np.meshgrid(xpoints, ypoints)


         # Plot the real and predicted surfaces.
         diffsurf = ax.plot_surface(xm, ym, np.vectorize(bestLasso.plotfunc)(xm, ym) - FrankeFunc

         # Customize the z axis.
         ax.set_zlim(-0.50, 0.5)
         ax.zaxis.set_major_locator(LinearLocator(10))
         ax.zaxis.set_major_formatter(FormatStrFormatter('%.02f'))

         # Add a color bar which maps values to colors.
         #fig.colorbar(realsurf, shrink=0.5, aspect=5)

         plt.show()

```




1.3 Parts d and e

1.3.1 Loading Terrain data (part d)

```
In [67]: #####
         # I got my terrain data by searching for blue mountains on earthexplorer
         # Used Blue Mountains, New South Wales, Australia -33.4100 150.3037
         #####

         # Load the terrain
         terrain1 = imread('s34_e150_1arc_v3.tif')
         # Show the terrain
         plt.figure()
         plt.title('The Blue Mountains')
         plt.imshow(terrain1, cmap='gray')
         plt.xlabel('X')
         plt.ylabel('Y')
         plt.show()

         terrain_array = np.array(terrain1)
         print(terrain_array.shape)
```



(3601, 3601)

1.4 Terrain analysis (part e) - model selection

```
In [68]: def terrain_function(row) :
          return terrain_array[(row[0], row[1])]

# Setup input data.
# We pick a 1000 (x,y) integer points at random
# the z's according to the fomula and with some noise
tnumberofpoints = 10000
tnoise = 0.5
tXY = np.random.randint(3600, size = (tnumberofpoints, 2))

tz = np.apply_along_axis(terrain_function, 1, tXY).reshape(-1,1)

# Now we make some "clean" data
# Models will never be trained on this, only tested against it
ctXY = np.random.randint(3600, size = (1000, 2))
ctz = np.apply_along_axis(terrain_function, 1, ctXY).reshape(-1,1)

print("we do cross fold validation to pick the best model")
cv_list = []
```

```

#first we try the OLS models

for n in range(1,6) :
    pOLS = polynomialOLS(n)
    cv = cross_validation(10, tXY, tz, pRidge)
    cv_list.append((cv,("OLS", n)))

# then we try Ridge
for n in range(1,6) :
    for a in [0.1, 0.2, 0.5, 1] :
        pRidge = polynomialRidge(n, a)
        cv = cross_validation(10, tXY, tz, pRidge)
        cv_list.append((cv,("Ridge", n, a)))

# then we try Ridge
for n in range(1,6) :
    for a in [0.1, 0.2, 0.5, 1] :
        pLasso = polynomialLasso(n, a)
        cv = cross_validation(10, tXY, tz, pLasso)
        cv_list.append((cv,("Lasso", n, a)))

cv_list.sort(key=lambda tup: tup[0]) #sort the list by the cv value

print(cv_list[0])

# Again the results are not exactly fixed, but when I ran it, I got
# Ridge with degree 3 with lambda 0.2

```

we do cross fold validation to pick the best model
(27589.21792591468, ('Ridge', 3, 0.5))

1.4.1 Working with chosen model

```

In [74]: # We will work with a degree 3 polynomial and a lambda of 0.5
# We test the actual statistics against the unseen (clean) data
tmodel = polynomialRidge(3, 0.5)
tmodel.train(tXY, tz)
print("The coefficients for our best model is:")
print(tmodel.coefficients())
test_against_new(ctXY, ctz, tmodel)

# Now we compute the bias and variance of our model
# and estimate the variance of the parameters beta

print("")
print("Using 9 folds to predict variance and bias")
b, v = BV_estimate(9, tXY, tz, tmodel)

```



```

print("The bias is %f" % b)
print("The variance is %f" % v)
print("")
print("Using 10 folds to compute variances of betas")
print(beta_variance(10, tXY, tz, tmodel))

```

The coefficients for our best model is:

```

[[ 7.74423356e+02]
 [ 2.63623576e-01]
 [-2.41755778e-01]
 [-6.09620512e-05]
 [-4.45552817e-05]
 [ 1.35688615e-05]
 [ 6.49809862e-09]
 [ 1.08423133e-08]
 [ 4.29164295e-09]
 [ 1.75566896e-10]]

```

Ridge regression, trying to fit a polynomial of degree 3

Lambda = 0.5000

MSE = 27383.656382

R2 score = 0.767798

Using 9 folds to predict variance and bias

The bias is 24636.566560

The variance is 5.116623

Using 10 folds to compute variances of betas

```

[1.42955179e+01 3.32090863e-06 1.19403923e-05 2.00656719e-13
 2.76789538e-13 5.22608739e-13 3.33498950e-19 8.82435959e-20
 2.44105857e-19 8.89518016e-20]

```

1.4.2 plotting terrain

```

In [76]: # Because of limited computing power of my laptop,
         # we will restrict our selves to the 500 by 500 top corner of the map
         # We train our model on some data from there,
         # Then plot it and the real map

```

```

restricted_XY = np.random.randint(499, size = (10000, 2))
restricted_z = np.apply_along_axis(terrain_function, 1, restricted_XY).reshape(-1,1)

```

```

rc_XY = np.random.randint(499, size = (1000, 2))
rc_z = np.apply_along_axis(terrain_function, 1, rc_XY).reshape(-1,1)

```

```

tmodel = polynomialRidge(3, 0.5)
tmodel.train(restricted_XY, restricted_z)

```

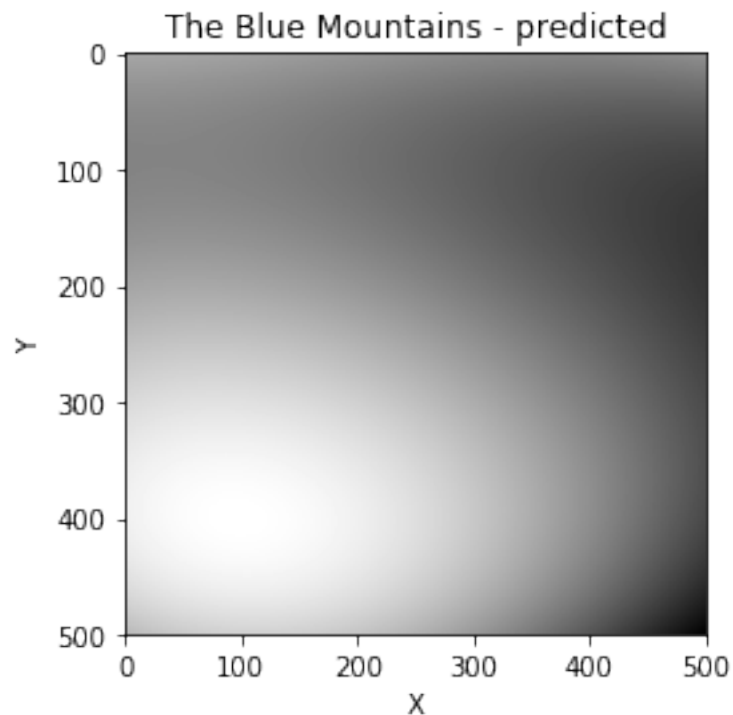
```

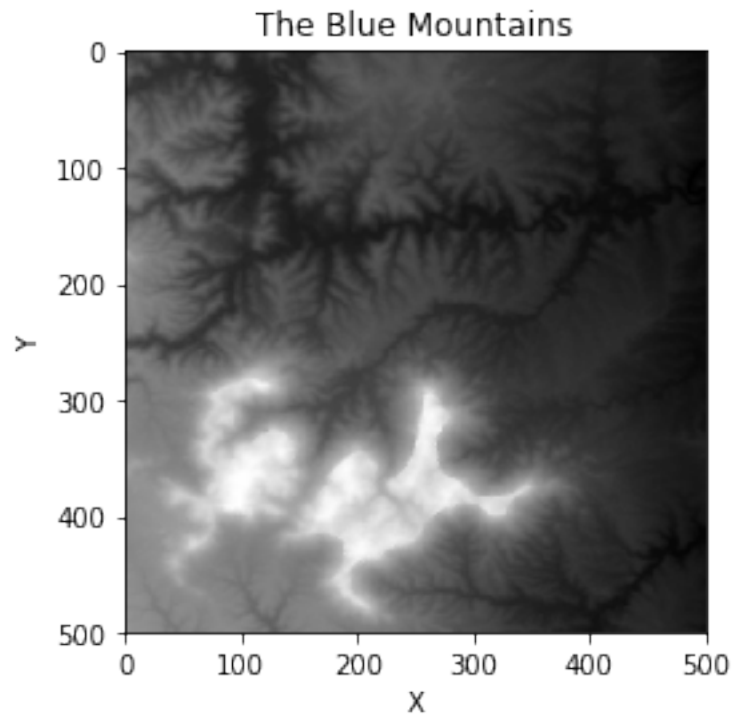
# plot a patch of our predicted landscape
my_terrain = np.rint(np.vstack([np.hstack([ tmodel.plotfunc(tx,ty) for ty in range(500)
plt.figure()
plt.title('The Blue Mountains - predicted')
plt.imshow(my_terrain, cmap='gray')
plt.xlabel('X')
plt.ylabel('Y')
plt.show()

# Show the terrain
plt.figure()
plt.title('The Blue Mountains')
plt.imshow(terrain_array[:500, :500], cmap='gray')
plt.xlabel('X')
plt.ylabel('Y')
plt.show()

print("The coefficients for our best model is:")
print(tmodel.coefficients())
test_against_new(rc_XY, rc_z, tmodel)

```





The coefficients for our best model is:

```
[[ 6.03684600e+02
  -2.08347169e+00
  -7.31907064e-03
   1.40890570e-02
   1.25217434e-03
  -1.45402890e-03
  -1.95642716e-05
   1.83078371e-06
  -7.47418512e-06
   2.55957023e-06]]
```

Ridge regression, trying to fit a polynomial of degree 3

Lambda = 0.5000

MSE = 6684.622435

R2 score = 0.713436
